

```
In [ ]: # Importing necessary Libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

In [ ]: # First, we observe the dataset and determine which dataset is relevant to our anal
hdb_df = pd.read_csv("C:\\Users\\KAVIPRIYA\\Desktop\\Resaleflatpricesbasedonregistr
hdb_df.head(10)
```

Out[]:

	month	town	flat_type	block	street_name	storey_range	floor_area_sqm	flat_model
0	2017-01	ANG MO KIO	2 ROOM	406	ANG MO KIO AVE 10	10 TO 12	44.0	Improved
1	2017-01	ANG MO KIO	3 ROOM	108	ANG MO KIO AVE 4	01 TO 03	67.0	New Generation
2	2017-01	ANG MO KIO	3 ROOM	602	ANG MO KIO AVE 5	01 TO 03	67.0	New Generation
3	2017-01	ANG MO KIO	3 ROOM	465	ANG MO KIO AVE 10	04 TO 06	68.0	New Generation
4	2017-01	ANG MO KIO	3 ROOM	601	ANG MO KIO AVE 5	01 TO 03	67.0	New Generation
5	2017-01	ANG MO KIO	3 ROOM	150	ANG MO KIO AVE 5	01 TO 03	68.0	New Generation
6	2017-01	ANG MO KIO	3 ROOM	447	ANG MO KIO AVE 10	04 TO 06	68.0	New Generation
7	2017-01	ANG MO KIO	3 ROOM	218	ANG MO KIO AVE 1	04 TO 06	67.0	New Generation
8	2017-01	ANG MO KIO	3 ROOM	447	ANG MO KIO AVE 10	04 TO 06	68.0	New Generation
9	2017-01	ANG MO KIO	3 ROOM	571	ANG MO KIO AVE 3	01 TO 03	67.0	New Generation

```
In [ ]: # In my analysis, I do not consider street name, block and flat model is relevant
hdb_df = hdb_df.drop(['month', 'street_name', 'flat_model', 'lease_commence_date', 'bl
```

```
In [ ]: # Let's rename the column so it will be clearer
hdb_df = hdb_df.rename(columns={'flat_type': 'number_of_rooms', 'storey_range': 'store
```

```
In [ ]: # I assume EXECUTIVE is equal to a 6 room (5 room + 1 study room). MULTI-GENERATION
hdb_df['number_of_rooms'] = hdb_df['number_of_rooms'].str.replace(r'EXECUTIVE', '6 R
hdb_df['number_of_rooms'] = hdb_df['number_of_rooms'].str.replace(r'MULTI-GENERATIO
hdb_df['number_of_rooms'] = hdb_df['number_of_rooms'].str.replace(r'ROOM', '', regex=
```

```
In [ ]: # I assume that rather we use floor range, I the possible highest floor within the
hdb_df['storey'] = hdb_df['storey'].str[-2:].astype('int')
```

```
In [ ]: # I revise the format of the data in the remaining lease to be quantifiable (change
hdb_df['remaining_lease'] = hdb_df['remaining_lease'].str.split(' ')
hdb_df['remaining_lease'] = hdb_df['remaining_lease'].apply(lambda x: (float(x[0]))+
hdb_df.head()
```

```
Out [ ]:
```

	town	number_of_rooms	storey	floor_area_sqm	remaining_lease	resale_price
0	ANG MO KIO	2	12	44.0	61.333333	232000.0
1	ANG MO KIO	3	3	67.0	60.583333	250000.0
2	ANG MO KIO	3	3	67.0	62.416667	262000.0
3	ANG MO KIO	3	6	68.0	62.083333	265000.0
4	ANG MO KIO	3	3	67.0	62.416667	265000.0

```
In [ ]: # Observe whether there is missing data or not.
hdb_df.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 165159 entries, 0 to 165158
Data columns (total 6 columns):
#   Column          Non-Null Count  Dtype
---  -
0   town             165159 non-null object
1   number_of_rooms  165159 non-null int32
2   storey           165159 non-null int32
3   floor_area_sqm   165159 non-null float64
4   remaining_lease  165159 non-null float64
5   resale_price     165159 non-null float64
dtypes: float64(3), int32(2), object(1)
memory usage: 6.3+ MB
```

```
In [ ]: hdb_df.isna().sum()
```

```
Out [ ]: town             0
number_of_rooms         0
storey                  0
floor_area_sqm          0
remaining_lease         0
resale_price            0
dtype: int64
```

```
In [ ]: # First, we want to see the distribution of HDB resale price in Singapore
sns.displot(hdb_df['resale_price'])
```

Out[]: <seaborn.axisgrid.FacetGrid at 0x27883ae6530>




In []: *# Let see the statistic information of the data*
 hdb_df.describe()

Out[]:

	number_of_rooms	storey	floor_area_sqm	remaining_lease	resale_price
count	165159.000000	165159.000000	165159.000000	165159.000000	1.651590e+05
mean	4.131510	9.768387	97.340378	74.703376	4.888170e+05
std	0.917013	5.948314	24.026982	13.821650	1.691271e+05
min	1.000000	3.000000	31.000000	42.250000	1.400000e+05
25%	3.000000	6.000000	82.000000	63.500000	3.650000e+05
50%	4.000000	9.000000	93.000000	74.666667	4.600000e+05
75%	5.000000	12.000000	112.000000	87.833333	5.800000e+05
max	6.000000	51.000000	249.000000	97.750000	1.500000e+06

In []: *# Let us see the relation between each parameters*
 sns.heatmap(hdb_df.describe(), cmap="YlGnBu")

Out[]: <Axes: >



```
In [ ]: hdb_df['town'].unique()
```

```
Out[ ]: array(['ANG MO KIO', 'BEDOK', 'BISHAN', 'BUKIT BATOK', 'BUKIT MERAH',
               'BUKIT PANJANG', 'BUKIT TIMAH', 'CENTRAL AREA', 'CHOA CHU KANG',
               'CLEMENTI', 'GEYLANG', 'HOUGANG', 'JURONG EAST', 'JURONG WEST',
               'KALLANG/WHAMPOA', 'MARINE PARADE', 'PASIR RIS', 'PUNGGOL',
               'QUEENSTOWN', 'SEMBAWANG', 'SENGKANG', 'SERANGOON', 'TAMPINES',
               'TOA PAYOH', 'WOODLANDS', 'YISHUN'], dtype=object)
```

```
In [ ]: hdb_df = hdb_df.replace(dict.fromkeys(['SEMBAWANG', 'SENGKANG', 'WOODLANDS', 'YISHUN']
hdb_df = hdb_df.replace(dict.fromkeys(['BUKIT MERAH', 'BUKIT TIMAH', 'QUEENSTOWN'], '
hdb_df = hdb_df.replace(dict.fromkeys(['BEDOK', 'GEYLANG', 'HOUGANG', 'KALLANG/WHAMPOA
hdb_df = hdb_df.replace(dict.fromkeys(['BUKIT BATOK', 'BUKIT PANJANG', 'CHOA CHU KANG
hdb_df = hdb_df.replace(dict.fromkeys(['ANG MO KIO', 'CENTRAL AREA', 'BISHAN', 'MARINE
```

```
In [ ]: # Let's check whether the data replacement was done properly
hdb_df['town'].unique()
```

```
Out[ ]: array(['CENTRAL', 'EAST', 'WEST', 'SOUTH', 'NORTH'], dtype=object)
```

```
In [ ]: hdb_df = hdb_df.rename(columns={'town': 'region'})
```

```
In [ ]: X = hdb_df.iloc[:, :-1].values
        y = hdb_df.iloc[:, -1].values
```

```
In [ ]: from sklearn.compose import ColumnTransformer
        from sklearn.preprocessing import OneHotEncoder
        ct = ColumnTransformer(transformers=[('encoder', OneHotEncoder(), [0])], remainder=
        X = np.array(ct.fit_transform(X))
```

```
In [ ]: X[1,:]
```

```
Out[ ]: array([1.0, 0.0, 0.0, 0.0, 0.0, 3, 3, 67.0, 60.583333333333336],
            dtype=object)
```

```
In [ ]: from sklearn.model_selection import train_test_split
        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2)
```

```
In [ ]: # Import necessary library to evaluate the performance of each machine Learning mo
        from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error
```

```
In [ ]: from sklearn.linear_model import LinearRegression
        mlr = LinearRegression()
        mlr.fit(X_train, y_train)
        mlr_ypred = mlr.predict(X_test)
        mlr_acc = r2_score(y_test, mlr_ypred)
        mlr_acc
```

```
Out[ ]: 0.6669831160774795
```

```
In [ ]: from sklearn.preprocessing import PolynomialFeatures
        from sklearn.linear_model import LinearRegression
        # Let's determine the best degree for polynomial
        for n in range(2,5):
            poly_reg = PolynomialFeatures(degree = n)
            X_poly = poly_reg.fit_transform(X_train)
            pr = LinearRegression()
            pr.fit(X_poly, y_train)
            poly_ypred = pr.predict(poly_reg.transform(X_test))
            poly_acc = r2_score(y_test, poly_ypred)
            poly_rmse = np.sqrt(mean_squared_error(y_test, poly_ypred))
            print(r'The accuracy of polynomial regression with degree of {} is {}'.format(n, poly_acc))
            print(r'The RMSE of polynomial regression with degree of {} is {}'.format(n, poly_rmse))
```

The accuracy of polynomial regression with degree of 2 is 0.7348388198462854
 The RMSE of polynomial regression with degree of 2 is 86941.23746903137
 The accuracy of polynomial regression with degree of 3 is 0.7475712228902434
 The RMSE of polynomial regression with degree of 3 is 84828.20503583294
 The accuracy of polynomial regression with degree of 4 is 0.7570782130310414
 The RMSE of polynomial regression with degree of 4 is 83215.47171211109

```
In [ ]: from sklearn.linear_model import Ridge
        ridge_r = Ridge()
        ridge_r.fit(X_train, y_train)
        ridge_ypred = ridge_r.predict(X_test)
```

```
ridge_acc = r2_score(y_test,ridge_ypred)
ridge_acc
```

Out[]: 0.6669710360456889

```
In [ ]: from sklearn.linear_model import Lasso
lasso_r = Lasso(max_iter=100000)
lasso_r.fit(X_train, y_train)
lasso_ypred = lasso_r.predict(X_test)
lasso_acc = r2_score(y_test,lasso_ypred)
lasso_acc
```

Out[]: 0.6669711175183721

```
In [ ]: from sklearn.linear_model import ElasticNet
EN_r = ElasticNet()
EN_r.fit(X_train, y_train)
EN_ypred = EN_r.predict(X_test)
EN_acc = r2_score(y_test,EN_ypred)
EN_acc
```

Out[]: 0.575702371290322

```
In [ ]: from sklearn.tree import DecisionTreeRegressor
tree_r = DecisionTreeRegressor()
tree_r.fit(X_train, y_train)
tree_ypred = tree_r.predict(X_test)
tree_acc = r2_score(y_test,tree_ypred)
tree_acc
```

Out[]: 0.7407413949535345

```
In [ ]: from sklearn.ensemble import RandomForestRegressor
forest_r = RandomForestRegressor(n_estimators = 10)
forest_r.fit(X_train, y_train)
forest_ypred = forest_r.predict(X_test)
forest_acc = r2_score(y_test,forest_ypred)
forest_acc
```

Out[]: 0.8037901498807709

```
In [ ]: # Accuracy score for multi linear regression
mlr_acc = r2_score(y_test,mlr_ypred)
mlr_rmse = np.sqrt(mean_squared_error(y_test,mlr_ypred))
# Evaluation for polynomial regression has been calculated in finding the best degr
# Evaluation for ridge regression
ridge_acc = r2_score(y_test,ridge_ypred)
ridge_rmse = np.sqrt(mean_squared_error(y_test,ridge_ypred))
# Evaluation for lasso regression
lasso_acc = r2_score(y_test,lasso_ypred)
lasso_rmse = np.sqrt(mean_squared_error(y_test,lasso_ypred))
# Evaluation for elastic net regression
EN_acc = r2_score(y_test,EN_ypred)
EN_rmse = np.sqrt(mean_squared_error(y_test,EN_ypred))
# Evaluation for decision trees regression
```

```

tree_acc = r2_score(y_test, tree_ypred)
tree_rmse = np.sqrt(mean_squared_error(y_test, tree_ypred))
# Evaluation for elastic random forest regression
forest_acc = r2_score(y_test, forest_ypred)
forest_rmse = np.sqrt(mean_squared_error(y_test, forest_ypred))
# Let's put it as a list and compare it in a bar chart
model_acc_score = [mlr_acc, poly_acc, ridge_acc, lasso_acc, EN_acc, tree_acc, forest_acc]
model_rmse = [mlr_rmse, poly_rmse, ridge_rmse, lasso_rmse, EN_rmse, tree_rmse, forest_rmse]
model_list = ['Multi Linear', 'Polynomial', 'Ridge', 'Lasso', 'Elastic Net', 'Decision Tree', 'Random Forest']
model_result_df = pd.DataFrame(
    {'Model': model_list,
     'Accuracy score': model_acc_score,
     'RMSE': model_rmse}
)

fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(15, 10))
sns.barplot(data=model_result_df, x='Model', y='Accuracy score', ax=ax1, order=model_list)
ax1 = ax1.set_xticklabels(ax1.get_xticklabels(), rotation=90)
sns.barplot(data=model_result_df, x='Model', y='RMSE', ax=ax2, order=model_list)
ax2 = ax2.set_xticklabels(ax2.get_xticklabels(), rotation=90)


```

C:\Users\KAVIPRIYA\AppData\Local\Temp\ipykernel_11388\1401320706.py:31: UserWarning: set_xticklabels() should only be used with a fixed number of ticks, i.e. after set_ticks() or using a FixedLocator.

ax1 = ax1.set_xticklabels(ax1.get_xticklabels(), rotation=90)

C:\Users\KAVIPRIYA\AppData\Local\Temp\ipykernel_11388\1401320706.py:33: UserWarning: set_xticklabels() should only be used with a fixed number of ticks, i.e. after set_ticks() or using a FixedLocator.

ax2 = ax2.set_xticklabels(ax2.get_xticklabels(), rotation=90)




```
In [ ]: # We know that after One Hot Encoding, the value of Central, East, North, South, and
area = ['Central', 'East', 'North', 'South', 'West']
pred_price_central = forest_r.predict([[1.0, 0.0, 0.0, 0.0, 0.0, 4, 9, 95.0, 75]])[0]
pred_price_east = forest_r.predict([[0.0, 1.0, 0.0, 0.0, 0.0, 4, 9, 95.0, 75]])[0]
pred_price_north = forest_r.predict([[0.0, 0.0, 1.0, 0.0, 0.0, 4, 9, 95.0, 75]])[0]
pred_price_south = forest_r.predict([[0.0, 0.0, 0.0, 1.0, 0.0, 4, 9, 95.0, 75]])[0]
pred_price_west = forest_r.predict([[0.0, 0.0, 0.0, 0.0, 1.0, 4, 9, 95.0, 75]])[0]
resale_price = [pred_price_central, pred_price_east, pred_price_north, pred_price_south, pred_price_west]
predict_df = pd.DataFrame(list(zip(area, resale_price)),
                           columns=['Area', 'Predicted HDB price (SGD)'])
predict_df.round()
```

```
Out[ ]:
```

	Area	Predicted HDB price (SGD)
0	Central	663200.0
1	East	743189.0
2	North	477300.0
3	South	776139.0
4	West	515900.0